

# SQL Injection Isn't Dead

## Smuggling Queries at the Protocol Level

**Paul Gerste – DEF CON 32 – August 10, 2024**

# *SQL INJECTION* **LOWER** **DECK#S**



HGETALL user:1

SELECT \* FROM users WHERE id=1

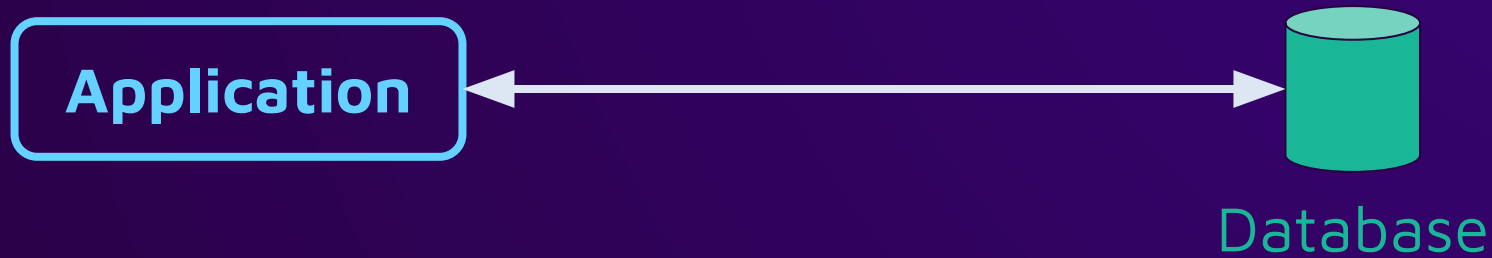
db.users.find({  
 id: 1,  
})



# Teaser

```
func getUser(w http.ResponseWriter, req *http.Request) (user User) {  
    body, _ := io.ReadAll(req.Body)  
    id := string(body)  
    db.QueryRow("SELECT * FROM users WHERE id=$1", id).Scan(&user)  
    // ...  
}
```

# Teaser



# Teaser





SELECT \* FROM speakers


name	role	company	team
Paul Gerste	Vuln Researcher	Sonar	R&D

(1 row)

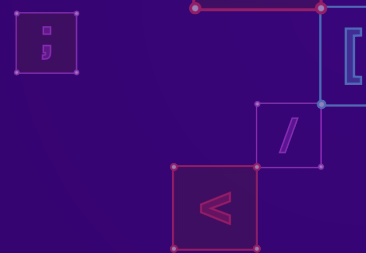


```
SELECT * FROM speakers INNER JOIN companies
```

name	role	company	team
Paul Gerste	Vuln Researcher	Sonar	R&D

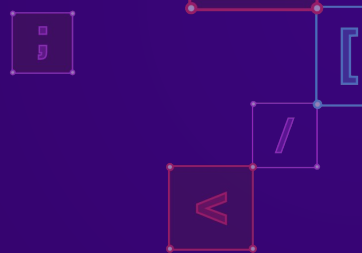
logo	name	description
	Sonar	The home of Clean Code

(1 row)



# Outline

- The Idea
- Attacking Database Wire Protocols
  - PostgreSQL
  - MongoDB
- Real-World Applicability
- Future Research
- Takeaways

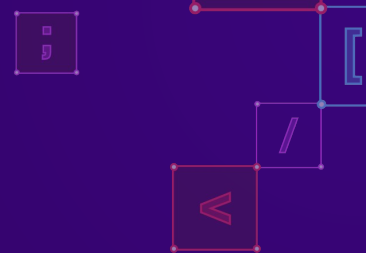


# The Idea

Request smuggling, but for binary protocols

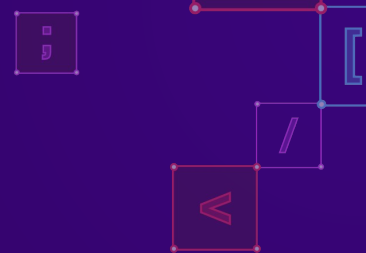
# Prior Art

- James Kettle: HTTP Desync Attacks
  - Cause disagreement over the end of HTTP requests
- Example root causes:
  - Text parsing: chunked vs. `[\t]chunked`
  - Logical: Content-Length vs. Transfer-Encoding
- What about other protocols?



# What About Binary Protocols?

- What are message boundaries here?
- Delimiters
  - E.g., null-terminated strings
- Length fields
  - E.g., Type-Length-Value (TLV) protocols



# Binary Protocols: Desync

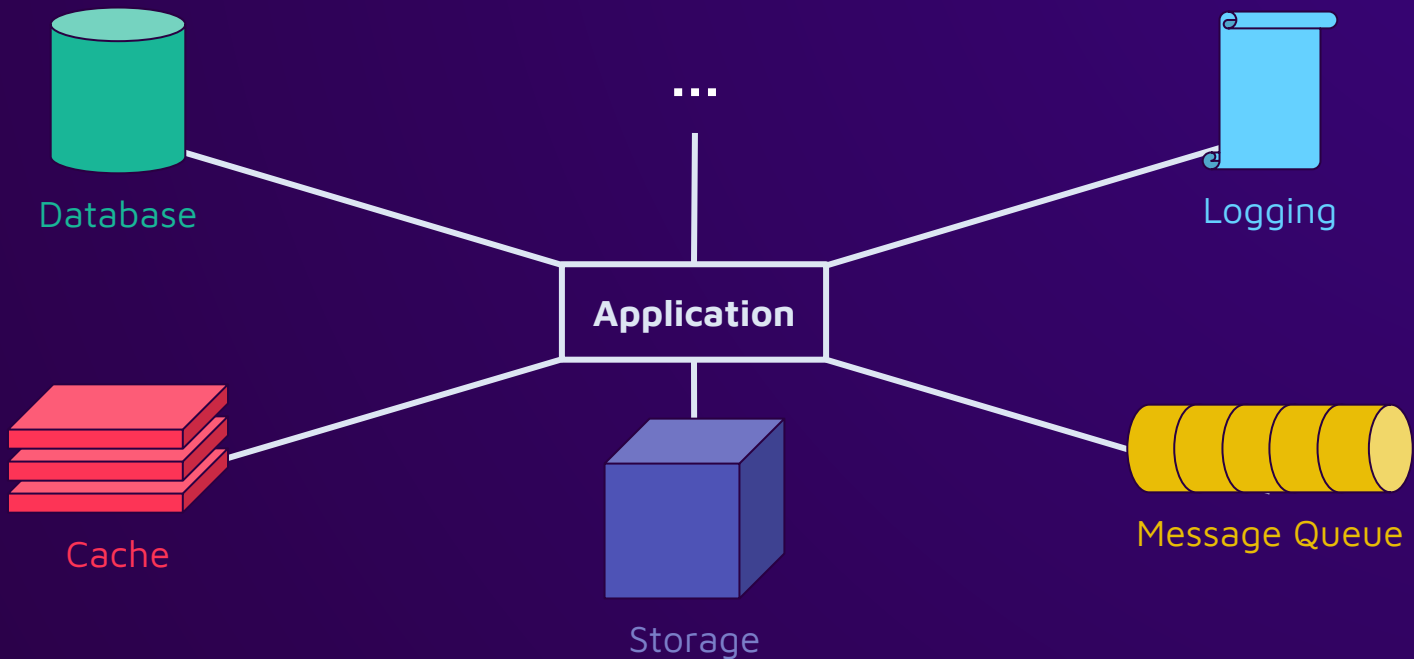
- Delimiters
  - Insert delimiters into values

# Binary Protocols: Desync

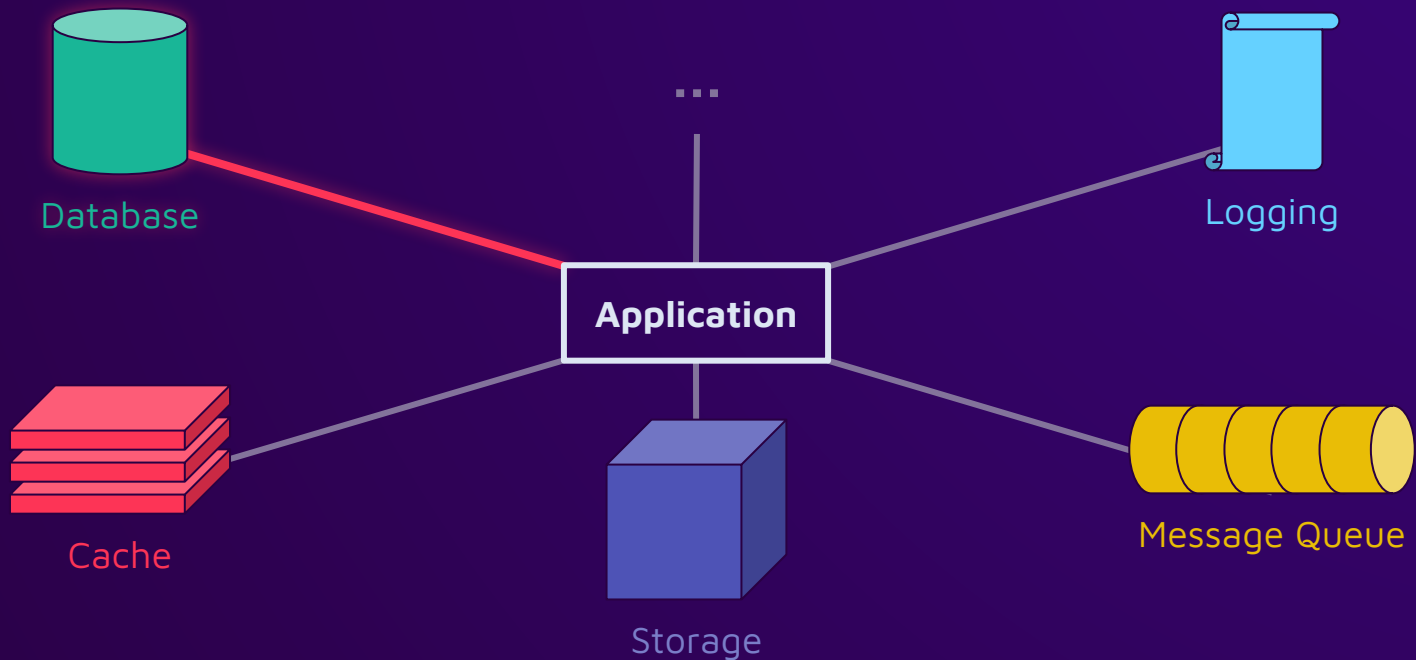
- Delimiters
  - Insert delimiters into values
- Length fields
  - 🤔
  - Endianness issues?
  - Overflows?

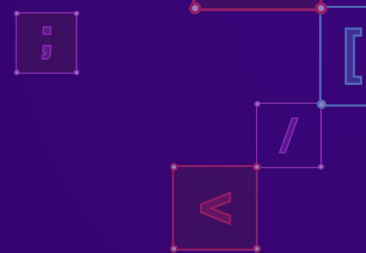


# Binary Protocols: Landscape



# Binary Protocols: Landscape

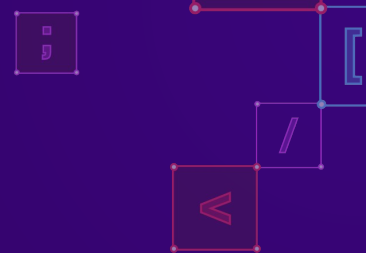




# Why Database Wire Protocols?

- Applicability
  - Almost every web app has a database
- Severity
  - Interesting data (e.g., PII)
  - Relevant data (e.g., for authentication)
- Exploitability
  - Most queries contain some user input

# Attacking Database Wire Protocols



# High-Level Protocol Comparison

- PostgreSQL
- MySQL
- Redis
- MongoDB

# High-Level Protocol Comparison

- **PostgreSQL**

Type	Length				Value...
'Q'	00	00	00	17	"SELECT ..."

- MySQL

- Redis

- MongoDB

# High-Level Protocol Comparison

- PostgreSQL

Type	Length				Value...
'Q'	00	00	00	17	"SELECT ..."

- **MySQL**

Length			Sequence	Value...
00	00	17	00	"SELECT ..."

- Redis

- MongoDB



# High-Level Protocol Comparison

- PostgreSQL

Type	Length				Value...
'Q'	00	00	00	17	"SELECT ..."

- MySQL

Length			Sequence	Value...
00	00	17	00	"SELECT ..."

- **Redis**

Type	Length	Delimiter	Value...	Delimiter
'+'	"17"	\r\n	"GET ..."	\r\n

- MongoDB

# High-Level Protocol Comparison

- PostgreSQL

Type	Length				Value...
'Q'	00	00	00	17	"SELECT ..."

- MySQL

Length			Sequence	Value...
00	00	17	00	"SELECT ..."

- Redis

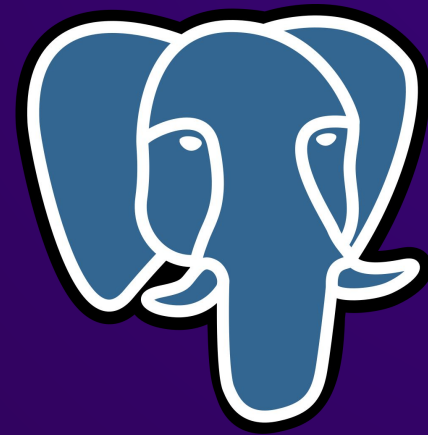
Type	Length	Delimiter	Value...	Delimiter
'+'	"17"	\r\n	"GET ..."	\r\n

- MongoDB

messageLength				requestID				responseTo			
17	00	00	00	00	00	00	00	00	00	00	00
opCode				value							
DD	07	00	00	...							

Case Study:

# PostgreSQL



# PostgreSQL Wire Protocol

Type	Length				Value...
'Q'	00	00	00	17	"SELECT ..."

- Type: 1-byte identifier
- Length: 4-byte integer
- Value

# PostgreSQL Wire Protocol

Type	Length				Value...
'Q'	00	00	00	17	"SELECT ..."

- Type: 1-byte identifier
- Length: 4-byte integer
- Value

Max value:  $2^{32}-1$

# PostgreSQL Wire Protocol

Type	Length				Value...
'Q'	00	00	00	17	"SELECT ..."

- Type: 1-byte identifier
- Length: 4-byte integer
- Value

Max value:  $2^{32}-1$



# The Bug: pgx

```
func (src *Bind) Encode(dst []byte) []byte {  
    dst = append(dst, 'B')  
    sp := len(dst)  
    // ...  
    pgio.SetInt32(dst[sp:], int32(len(dst[sp:])))  
    return dst  
}
```



# The Bug: pgx

```
func (src *Bind) Encode(dst []byte) []byte {  
    dst = append(dst, 'B')  
    sp := len(dst)  
    // ...  
    pgio.SetInt32(dst[sp:], int32(len(dst[sp:])))  
    return dst  
}
```

● Write message type

# The Bug: pgx

```
func (src *Bind) Encode(dst []byte) []byte {  
    dst = append(dst, 'B')  
    sp := len(dst) —● Save size offset  
    // ...  
    pgio.SetInt32(dst[sp:], int32(len(dst[sp:])))  
    return dst  
}
```

# The Bug: pgx

```
func (src *Bind) Encode(dst []byte) []byte {  
    dst = append(dst, 'B')  
    sp := len(dst)  
    // ...  
    pgio.SetInt32(dst[sp:], int32(len(dst[sp:])))  
    return dst  
}
```

● Build the rest

# The Bug: pgx

```
func (src *Bind) Encode(dst []byte) []byte {  
    dst = append(dst, 'B')  
    sp := len(dst)  
    // ...  
    pgio.SetInt32(dst[sp:], int32(len(dst[sp:])))  
    return dst  
}
```

Write size

# The Bug: pgx

```
func (src *Bind) Encode(dst []byte) []byte {  
    dst = append(dst, 'B')  
    sp := len(dst)  
    // ...  
    pgio.SetInt32(dst[sp:], int32(len(dst[sp:])))  
    return dst  
}
```

The message buffer

# The Bug: pgx

```
func (src *Bind) Encode(dst []byte) []byte {  
    dst = append(dst, 'B')  
    sp := len(dst)  
    // ...  
    pgio.SetInt32(dst[sp:], int32(len(dst[sp:])))  
    return dst  
}
```

Buffer length (int)

# The Bug: pgx

```
func (src *Bind) Encode(dst []byte) []byte {  
    dst = append(dst, 'B')  
    sp := len(dst)  
    // ...  
    pgio.SetInt32(dst[sp:], int32(len(dst[sp:])))  
    return dst  
}
```

Truncate to int32



# Message Size Overflow

Message 1					
Type	Length				Value
'Q'	00	00	00	08	"AAAA"

Size: 8 = 0x00000008

4 bytes length + 4 bytes data

Payload: "A" \* 4

# Message Size Overflow

Message 1					
Type	Length				Value
'Q'	FF	FF	FF	FF	"AA..."

Size:  $2^{32}-1 = 0x\text{FFFFFFFF}$

4 bytes length +  $2^{32}-5$  bytes data

Payload: "A" \* ( $2^{32} - 5$ )

# Message Size Overflow

Message 1					?					
Type	Length				Value	?	?			
'Q'	00	00	00	04	""	'A'	'A'	'A'		

Size:  $2^{32} + 4 = 0x1000000004$

4 bytes length +  $2^{32}$  bytes data

Payload: "A" \* (2\*\*32)

# Message Size Overflow

Message 1					Injected Message			
Type	Length				Value	Type	Length	
'Q'	00	00	00	04	""	'Q'	00	00

Size:  $2^{32} + 4 = 0x1000000004$

4 bytes length +  $2^{32}$  bytes data

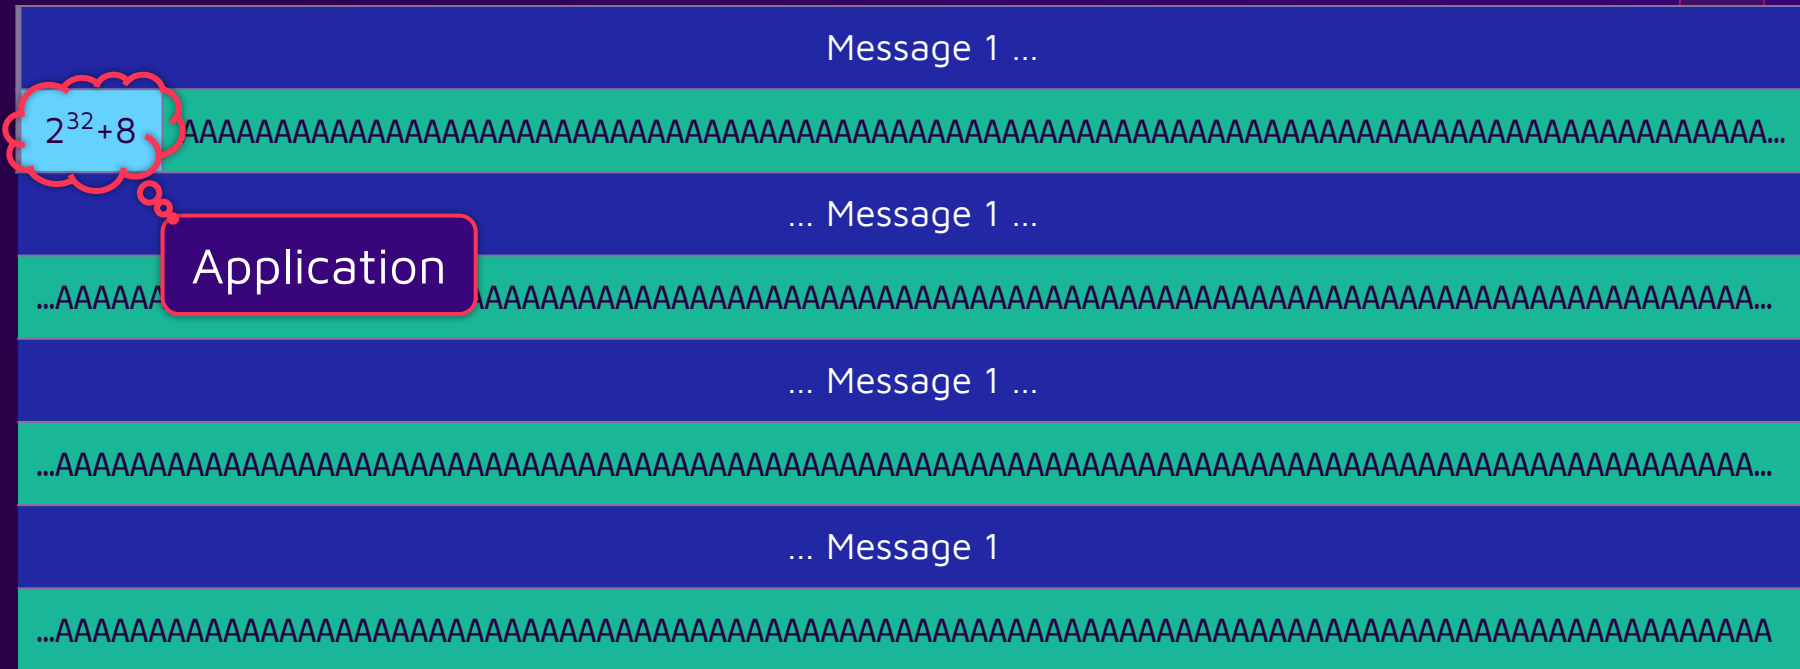
Payload: `fakeMsg + "A" * (2**32 - len(fakeMsg))`

# Message Size Overflow - Zoomed Out

Message 1	
8	AAAA



# Message Size Overflow - Zoomed Out





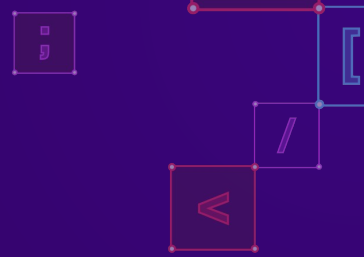


[illegible]

# Impact

- Inject entire SQL statements
  - Not limited to UNION, subqueries, etc.
  - Like stacked queries
- Read/write/delete all data in the DB
- Direct exfiltration is inconvenient
  - Application only processes the first DB response

# How does it look in the real world?



# How does it look in the real world?

```
id := "5831bfef"
```

```
conn.QueryRow("SELECT * FROM users WHERE id = $1", id)
```

Type	Length				Value
'Q'	00	00	00	2e	SELECT * FROM users WHERE id = '5831bfef'\x00

# How does it look in the real world?

```
id := strings.Repeat("A", 1<<32)
```

```
conn.QueryRow("SELECT * FROM users WHERE id = $1", id)
```

Type	Length				Value
'Q'	00	00	00	26	SELECT * FROM users WHERE id = 'AAAAAAAAAAAAAAAAAAAA...

0x26 = 38

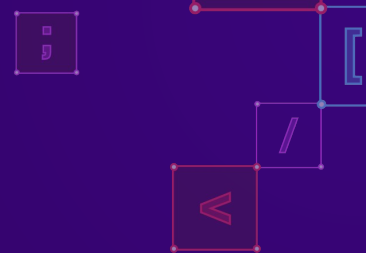
# How does it look in the real world?

```
id := strings.Repeat("A", 1<<32)
```

```
conn.QueryRow("SELECT * FROM users WHERE id = $1", id)
```

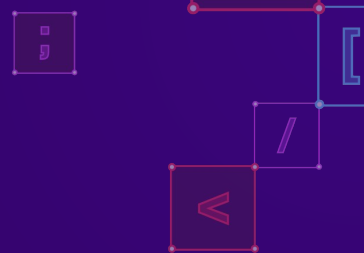
Type	Length				Value	Type	Length		
'Q'	00	00	00	29	SELECT * FROM users WHERE id = 'A'	'Q'	00		

How to know this offset? —



# Crafting a Payload

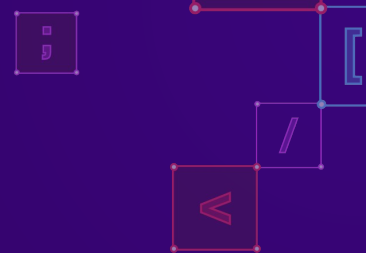
- Offset depends on the query
  - Where is the injection point?
  - How long is the query?
- Calculate the offset when query is known
- What if it's not?



# Crafting a Payload

- Naïve solution: Try all the offsets!
  - Need to send 4GB for each try
  - Takes time, creates noise
  - Risk of DoS
- Can we make it more reliable?





# Crafting a Payload: NOP Sled

- Idea: NOP sled
  - Use a lot of small messages
  - Hit start of a message → success
  - Hit something else → connection closed

# Crafting a Payload: NOP Sled

Smallest possible message

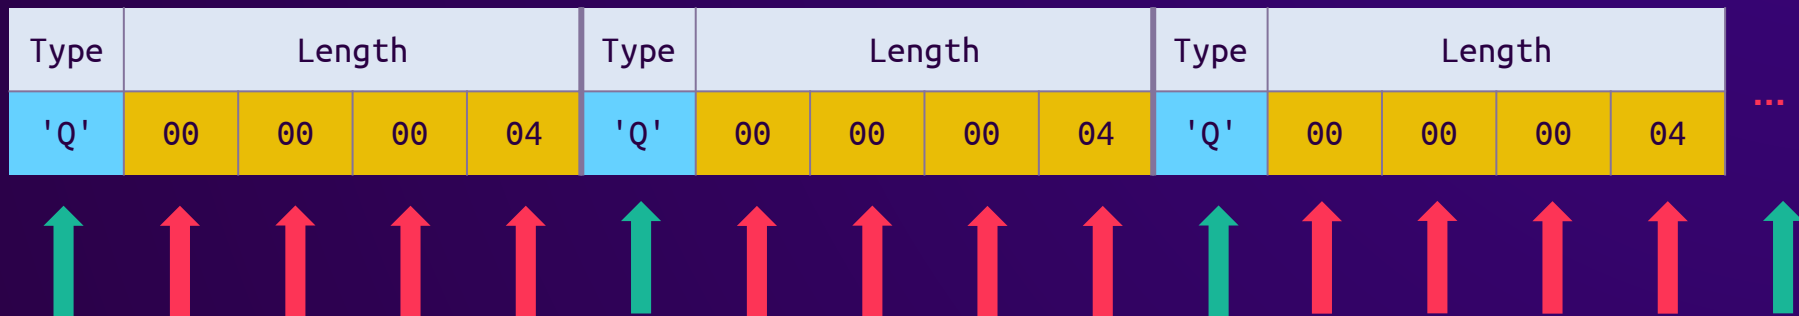
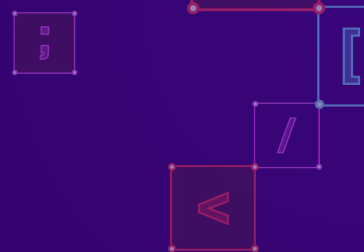
Type	Length			
'Q'	00	00	00	04

# Crafting a Payload: NOP Sled

Type	Length				Type	Length				Type	Length				...
'Q'	00	00	00	04	'Q'	00	00	00	04	'Q'	00	00	00	04	

Type	Length				Type	Length				Value
'Q'	00	00	00	04	'Q'	00	00	00	3B	INSERT INTO admins VALUES ...

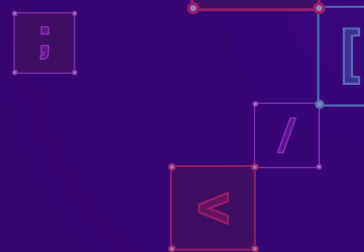
# Crafting a Payload: NOP Sled



# Crafting a Payload: NOP Sled

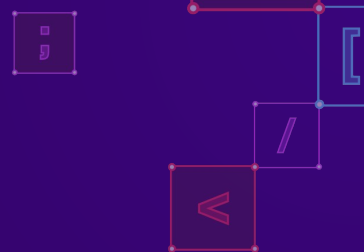
Pad	Type	Length				Type	Length				Type	Length			
A	'Q'	00	00	00	04	'Q'	00	00	00	04	'Q'	00	00	00	04
↑	↑	↑	↑	↑	↑	↑	↑	↑	↑	↑	↑	↑	↑	↑	↑

# Crafting a Payload: NOP Sled



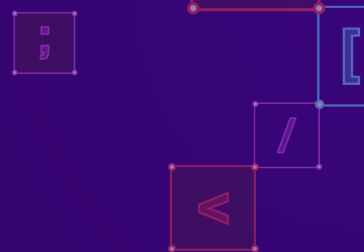
Pad		Type	Length				Type	Length				Type	Length			
A	A	'Q'	00	00	00	04	'Q'	00	00	00	04	'Q'	00	00	00	
↑	↑	↑	↑	↑	↑	↑	↑	↑	↑	↑	↑	↑	↑	↑	↑	↑

# Crafting a Payload: NOP Sled



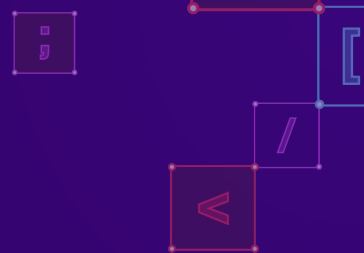
Pad			Type	Length				Type	Length				Type	Length		
A	A	A	'Q'	00	00	00	04	'Q'	00	00	00	04	'Q'	00	00	
↑	↑	↑	↑	↑	↑	↑	↑	↑	↑	↑	↑	↑	↑	↑	↑	

# Crafting a Payload: NOP Sled



Pad				Type	Length				Type	Length				Type	
A	A	A	A	'Q'	00	00	00	04	'Q'	00	00	00	04	'Q'	00
↑	↑	↑	↑	↑	↑	↑	↑	↑	↑	↑	↑	↑	↑	↑	↑





# Crafting a Payload: NOP Sled

- Success after  $\leq 5$  attempts!
  - 20% chance of success
  - Attack is repeatable, just change the offset
- Still have to send  $5 \times 4$  GB in the worst case
  - Can we make it even better?

# Crafting a Payload: Trampolines

- Can length bytes be valid types?

Type	Length			
'Q'	00	00	00	04



# Crafting a Payload: Trampolines

- Can length bytes be valid types?
  - Trampolines!

Type	Length			
'Q'	'Q'	'Q'	'Q'	'Q'
	?	?	?	?

# Crafting a Payload: Trampolines

- Can length bytes be valid types?
  - Trampolines!

Type	Length			
'Q'	51	51	51	51
	?	?	?	?

Q	Q	Q	Q	Q	S	S	S	S	S	B	B	B	B	B	E	E	E	E	E	Z	Z	Z	Z	Z	...	Q	?	?	?	?	Q	?	?	?	?	Q	?	?
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	-----	---	---	---	---	---	---	---	---	---	---	---	---	---

# Crafting a Payload: Trampolines

- Can length bytes be valid types?
  - Trampolines!

Type	Length			
'Q'	51	51	51	51
	?	?	?	?



Q	Q	Q	Q	Q	S	S	S	S	S	B	B	B	B	B	E	E	E	E	E	Z	Z	Z	Z	Z	...	Q	?	?	?	?	Q	?	?	?	?	Q	?	?
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	-----	---	---	---	---	---	---	---	---	---	---	---	---	---

# Crafting a Payload: Trampolines

- Can length bytes be valid types?
  - Trampolines!

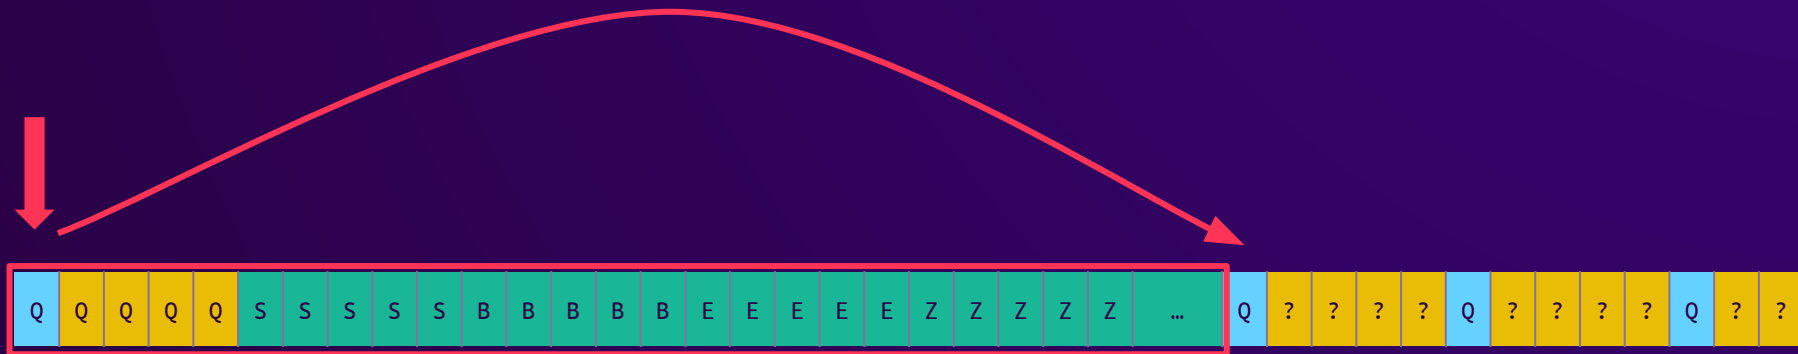
Type	Length			
'Q'	51	51	51	51
	?	?	?	?



# Crafting a Payload: Trampolines

- Can length bytes be valid types?
  - Trampolines!

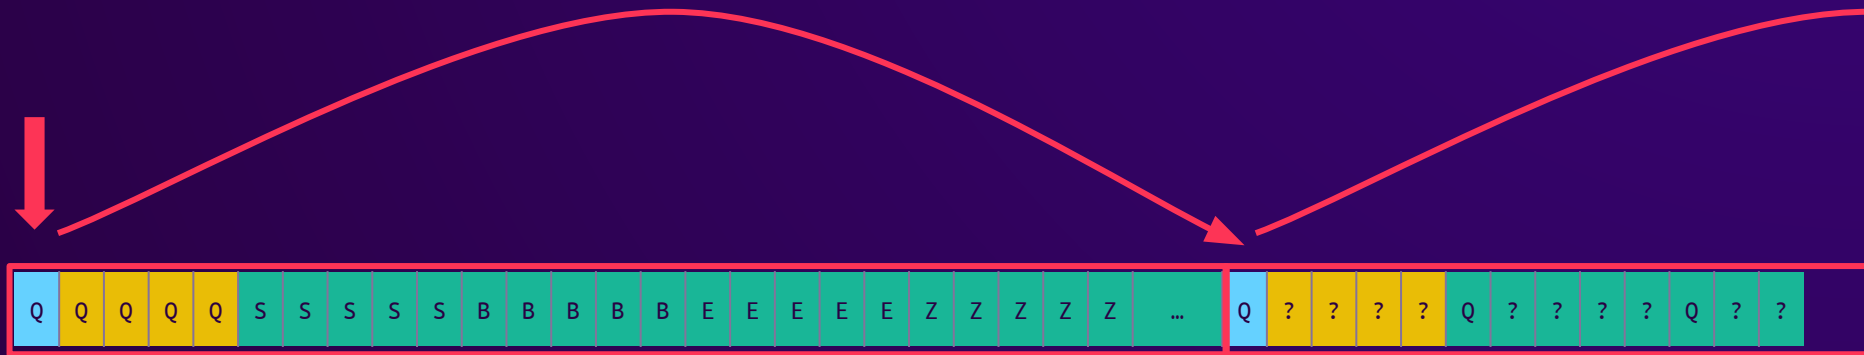
Type	Length			
'Q'	51	51	51	51
	?	?	?	?



# Crafting a Payload: Trampolines

- Can length bytes be valid types?
  - Trampolines!

Type	Length			
'Q'	51	51	51	51
	?	?	?	?





# Crafting a Payload: Trampolines

- Can length bytes be valid types?
  - Trampolines!

Type	Length			
'Q'	51	51	51	51
	?	?	?	?



Q	Q	Q	Q	Q	S	S	S	S	S	B	B	B	B	B	E	E	E	E	E	Z	Z	Z	Z	Z	...	Q	?	?	?	?	Q	?	?	?	?	Q	?	?
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	-----	---	---	---	---	---	---	---	---	---	---	---	---	---

# Crafting a Payload: Trampolines

- Can length bytes be valid types?
  - Trampolines!

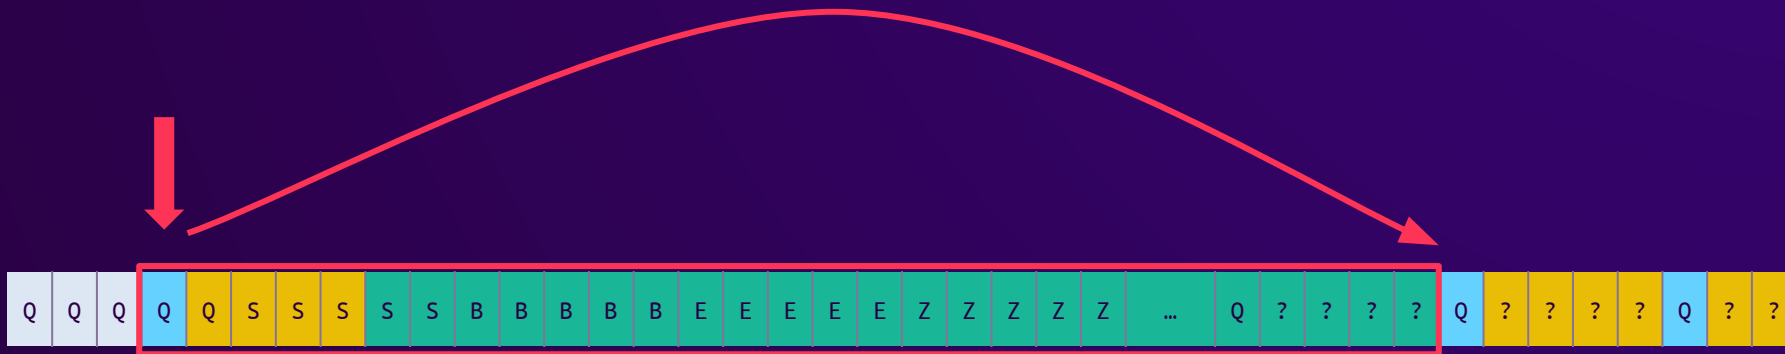
Type	Length			
'Q'	51	51	51	51
	?	?	?	?



# Crafting a Payload: Trampolines

- Can length bytes be valid types?
  - Trampolines!

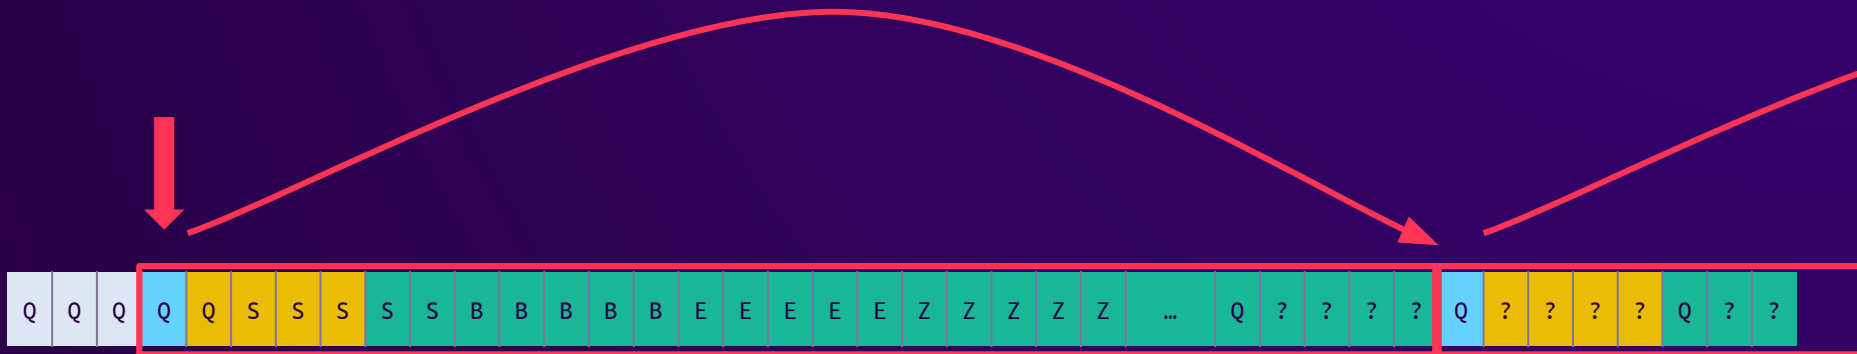
Type	Length			
'Q'	51	51	51	51
	?	?	?	?



# Crafting a Payload: Trampolines

- Can length bytes be valid types?
  - Trampolines!

Type	Length			
'Q'	51	51	51	51
	?	?	?	?



# Crafting a Payload: Trampolines

- Can length bytes be valid types?
  - Trampolines!

Type	Length			
'Q'	51	51	51	51
	?	?	?	?

# Crafting a Payload: Trampolines

- Can length bytes be valid types?
  - Trampolines!
- Max. logical size: 0x3fffffff
  - First size byte cannot be > 0x3f

Type	Length			
'Q'	51	51	51	51
				

# Crafting a Payload: Trampolines

- Can length bytes be valid types?
  - Trampolines!
- Max. logical size: 0x3fffffffff
  - First size byte cannot be > 0x3f

Type	Length				
3f	3f	3f	3f	3f	3f
?	?	?	?	?	?

# Crafting a Payload: Trampolines

- Can length bytes be valid types?
  - Trampolines!
- Max. logical size:  $0x3fffffffff$ 
  - First size byte cannot be  $> 0x3f$
- No valid message type  $\leq 0x3f$

Type	Length			
3f	3f	3f	3f	3f
✗	✓	✓	✓	✓



# Crafting a Payload: Trampolines

- Can length bytes be valid types?
  - Trampolines!
- Max. logical size: 0x3fffffff
  - First size byte cannot be > 0x3f
- No valid message type  $\leq 0x3f$
- Solution: alternating pattern

Type	Length			
'Q'	00	'Q'	00	'Q'
✓	✗	✓	✗	✓

# Crafting a Payload: Trampolines

- Every **2nd** byte is a valid type
  - Hit a valid type byte → success
  - Hit other bytes → connection closed
- Success after  $\leq 2$  attempts!
  - 50% chance of success
  - Attack is repeatable, just change the offset

Type	Length			
'Q'	00	'Q'	00	'Q'
✓	✗	✓	✗	✓

# Vulnerable Libraries

Language	Library	Vulnerable?	Exploitable?	Fixed Versions
Go	pgx	✓	✓	4.18.2, 5.5.4
	pg	✓	✓	none
	pgdriver	✓	✓	none
	pq	✓	✓	none
C#/.NET	Npgsql	✓	✓	4.0.14, 4.1.13, 5.0.18, 6.0.11, 7.0.7, 8.0.3
Java	pgjdbc	✗	✗	-
	pgjdbc-ng	✓	✗	-
	r2dbc-postgresql	✓	✗	-
JS/TS	pg	✓	✗	-
	pg-promise	✗	✗	-
	pogi	✓	✗	-
	postgres	✓	✗	-
	@vercel/postgres	✓	✗	-

# Disclosure Timeline

- Sent advisories in February 2024
- `pgx` fixed in March
- `Npgsql` fixed in May
- `pg` and `pgdriver` maintainer initially responded but then stopped
- `pq` maintainers never responded to issue/PR

# Exploitable Applications

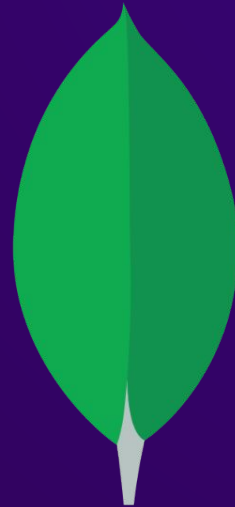


# Demo: Harbor

- Container registry
  - CNCF Graduate project
  - Part of VMware Tanzu Kubernetes
- Default configuration was vulnerable
- No authentication required
- Fixed in 2.11.0 by updating pgx <sup>[1]</sup>



# Case Study: MongoDB



# MongoDB Wire Protocol

messageLength				requestID				responseTo			
17	00	00	00	00	00	00	00	00	00	00	00
opCode				value							
DD	07	00	00	...							

- 4-byte length field
- Queries are BSON documents
  - Hierarchical objects
  - Serialized to TLV sections



# The Bug: mongodb

```
async fn write_to<T: AsyncWrite + Send + Unpin>(&self, mut writer: T) -> Result<()> {
    let sections = self.get_sections_bytes();
    let total_length = Header::LENGTH
        + std::mem::size_of::<u32>()
        + sections.len()
        + /* ... */;
    let header = Header {
        length: total_length as i32,
        // ...
    };
    header.write_to(&mut writer).await?;
    writer.write_u32_le(self.flags.bits()).await?;
    writer.write_all(&sections).await?;
    // ...
}
```

# The Bug: mongodb

```
async fn write_to<T: AsyncWrite + Send + Unpin>(&self, mut writer: T) -> Result<()> {
```

```
    let sections = self.get_sections_bytes();
```

```
    let total_length = Header::LENGTH
```

```
        + std::mem::size_of::<u32>()
```

```
        + sections.len()
```

```
        + /* ... */;
```

```
    let header = Header {
```

```
        length: total_length as i32,
```

```
        // ...
```

```
    };
```

```
    header.write_to(&mut writer).await?;
```

```
    writer.write_u32_le(self.flags.bits()).await?;
```

```
    writer.write_all(&sections).await?;
```

```
    // ...
```

```
}
```

• Get content bytes

# The Bug: mongodb

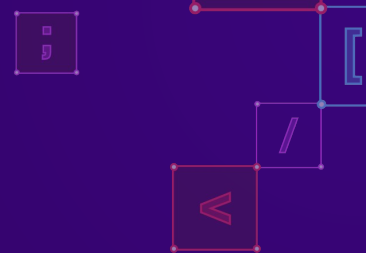
```
async fn write_to<T: AsyncWrite + Send + Unpin>(&self, mut writer: T) -> Result<()> {  
    let sections = self.get_sections_bytes();  
    let total_length = Header::LENGTH  
        + std::mem::size_of::<u32>()  
        + sections.len()  
        + /* ... */;  
    let header = Header {  
        length: total_length as i32,  
        // ...  
    };  
    header.write_to(&mut writer).await?;  
    writer.write_u32_le(self.flags.bits()).await?;  
    writer.write_all(&sections).await?;  
    // ...  
}
```

• Calculate message size (usize)

# The Bug: mongodb

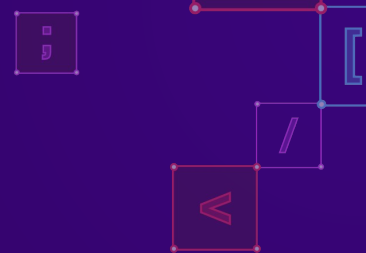
```
async fn write_to<T: AsyncWrite + Send + Unpin>(&self, mut writer: T) -> Result<()> {  
    let sections = self.get_sections_bytes();  
    let total_length = Header::LENGTH  
        + std::mem::size_of::<u32>()  
        + sections.len()  
        + /* ... */;  
    let header = Header {  
        length: total_length as i32,  
        // ...  
    };  
    header.write_to(&mut writer).await?;  
    writer.write_u32_le(self.flags.bits()).await?;  
    writer.write_all(&sections).await?;  
    // ...  
}
```

• Truncate to i32



# Crafting a Payload

- Avoid bad bytes
  - Payload must be valid UTF-8
- Problem:
  - Message type (dd 07) is already invalid
  - Size fields can become invalid



# Crafting a Payload

- Avoid bad bytes
  - Payload must be valid UTF-8
- Problem:
  - Message type (dd 07) is already invalid
  - Size fields can become invalid
- Solution:
  - Use metadata to create those bytes!

# Crafting a Payload

Query:

```
{  
  title: "The Wrath of Khan",  
  genre: "SciFi",  
  description: "...",  
}
```

BSON Document:

4800	0000	02	74	6974	6c65	0012	0000	0054
6865	2057	7261	7468	206f	6620	4b68	616e	
0002	6765	6e72	6500	0600	0000	5363	6946	
6900	0264	6573	6372	6970	7469	6f6e	0004	
0000	002e	2e2e	0000					

H...	title....	T
he	Wrath of Khan	
..	genre.....	SciF
i..	description..	
....	....	

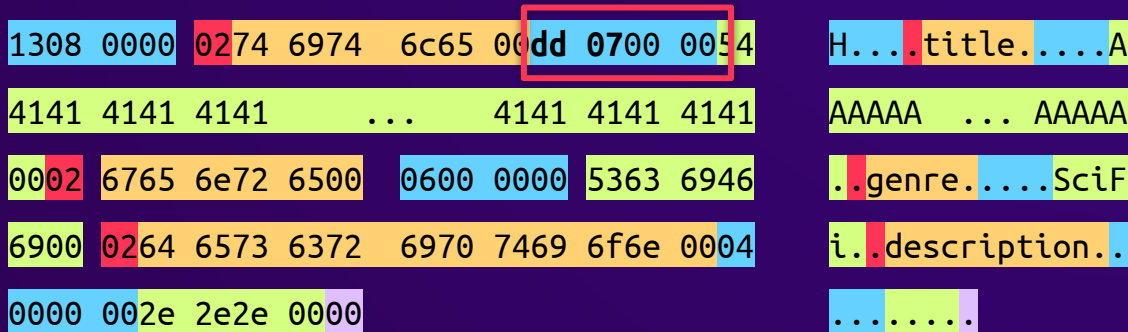
Length Type Key Value Other

# Crafting a Payload

Query:

```
{  
  title: "A" * (0x7dd - 1),  
  genre: "SciFi",  
  description: "...",  
}
```

BSON Document:



Length Type Key Value Other



# Vulnerable Libraries

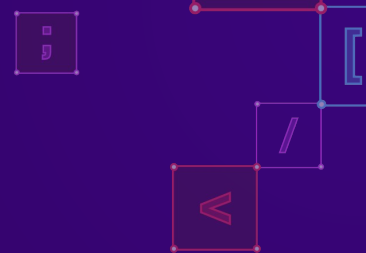
Language	Library	Vulnerable?	Exploitable?	Fixed Version
Rust	mongodb	✓	✓	2.8.2
Python	pymongo	✗	✗	-
Go	mongo	✗	✗	-
Java	mongo-java-driver	✗	✗	-
JavaScript	mongodb	✗	✗	-

- Sent advisory in February 2024
- mongodb fixed in March

# Real-World Applicability

# Constraints





# How Web Apps Handle Large Payloads

- Aren't apps limiting input sizes?
- Common protections:
  - Default body size limits
  - Maximum JSON/form decode sizes
  - Size-limiting reverse proxies
  - ... and more

# How Web Apps Handle Large Payloads

- Potential bypasses
  - Unprotected endpoints
  - Compression
  - WebSockets
  - Alternate body types
  - Server-side creation

# How Web Apps Handle Large Payloads

- Potential bypasses
  - **Unprotected endpoints**
  - Compression
  - WebSockets
  - Alternate body types
  - Server-side creation
- Some have no default limits
- Some explicitly disable limits
  - Harbor

# How Web Apps Handle Large Payloads

- Potential bypasses
  - Unprotected endpoints
  - **Compression**
  - WebSockets
  - Alternate body types
  - Server-side creation

- Some enforce size limits **before** decompression
  - Nginx
  - Fastify

# How Web Apps Handle Large Payloads

- Potential bypasses
  - Unprotected endpoints
  - Compression
  - **WebSockets**
  - Alternate body types
  - Server-side creation
- Compression support
- Large message size
- Many filters don't apply

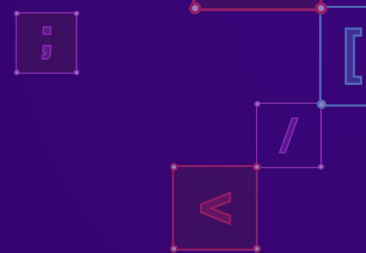


# How Web Apps Handle Large Payloads

- Potential bypasses
  - Unprotected endpoints
  - Compression
  - WebSockets
  - **Alternate body types**
  - Server-side creation
- Some filters don't apply
  - E.g., multipart forms

# How Web Apps Handle Large Payloads

- Potential bypasses
  - Unprotected endpoints
  - Compression
  - WebSockets
  - Alternate body types
  - **Server-side creation**
- Create strings on the server side
  - SSRF, templates, i18n, etc.
- Can depend on business logic



# Language Comparison

- How well do languages handle big payloads?
  - How big can strings/buffers be?
- Are integer overflows silent?

# Language Comparison: Large Payloads

Language	Max. String Size	Max. Buffer Size
Go	$> 2^{32}$	$> 2^{32}$
Java	$2^{31}-1$	$2^{31}-1$
C#	$2^{31}-1$	$> 2^{32}$
JS	$2^{29}-24 *$	$> 2^{32} *$
Python	$> 2^{32}$	$> 2^{32}$
Rust	$> 2^{32}$	$> 2^{32}$

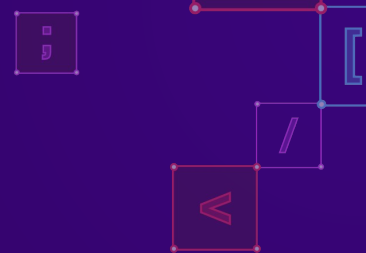
Only considering 64-bit versions.

\* Depends on the implementation

# Language Comparison: Integer Overflows

Language	Silent Addition Overflow?	Silent Serialization Overflow?
Go	Yes	N/A *
Java	Yes	N/A *
C#	Yes	N/A *
JS	No	Depends on impl.
Python	No	No
Rust	In release builds	N/A *

\* Type system prevents overflows. Devs have to check for overflows, leading to bugs



# Real-World Applicability

- Can I send large payloads?
  - A lot of times, yes!
- Can integers silently overflow/truncate?
  - In many languages, yes!
- Can I exploit real-world apps with this?
  - Absolutely!

# Future Research

# Safety First: No DoS Please!



Do not send large payloads to third-party systems!

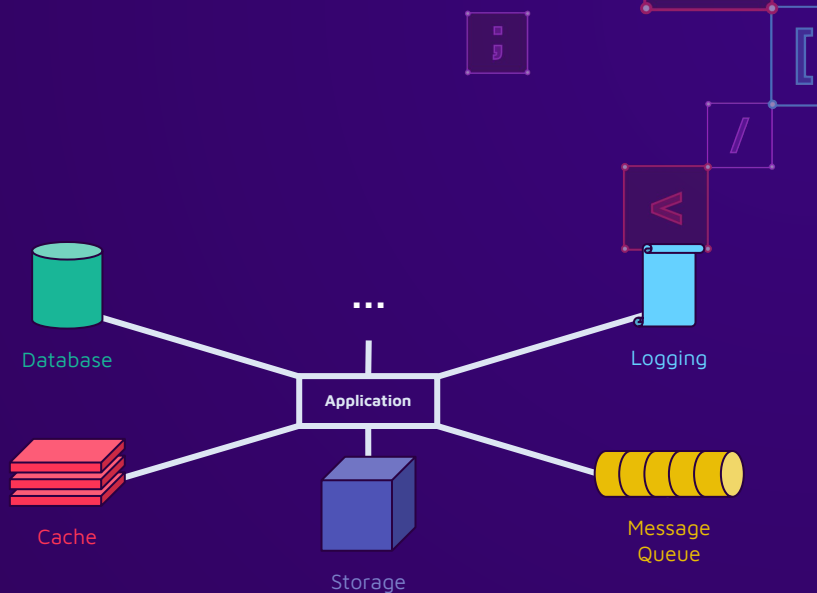


## Non-Invasive Detection

- White-box tests are harmless
  - Just set up your own test environment
- How to test this black-box?
  - Sending large payloads risks DoS
- More research and tools needed!
  - Can we safely detect vulnerable libraries?
  - Build tools to test this safely

# Research More!

- More protocols
  - Other databases
  - Caches, message queues, ...
- Find more desync techniques
  - What about delimiters?
- More "large payload" methods
  - New ways to bypass limits
  - Generic server-side creation techniques



# Research More!

- All this was about **4-byte** length fields
- What about 2-byte fields?
  - Much easier to exploit (65KB vs. 4GB)
  - More to come in the future 🙄🙄

# Conclusion

# Takeaways

- Integer overflows are still relevant in memory-safe languages
- Sending large amounts of data is feasible
- SQL injection isn't dead
  - If you can't hack it, just go a level deeper!

# Thank you!



@Sonar\_Research



@SonarResearch@infosec.exchange



<https://sonarsource.com>



@pspaul95



@pspaul@infosec.exchange